



Services

Performance. Scale. Reliability — Engineered.

Mission

KPI99 helps organizations operating mission-critical systems identify and eliminate performance and capacity limits before they impact customers, service levels, or budgets. Our focus is on measurable outcomes — reduced latency, improved capacity clarity, and optimized infrastructure efficiency.

Services Overview

1. Performance Health Audit

Purpose:

Comprehensive assessment of current performance and capacity posture.

Duration: 2–3 weeks

Investment: \$10,000

Scope

- JVM GC, heap, thread, and memory analysis
- Infrastructure utilization (CPU, memory, disk, network)
- Load profiles and peak demand behavior
- Bottleneck detection and risk identification
- Cloud and infrastructure inefficiency review

Deliverables

- Executive summary (non-technical)
- Detailed performance findings
- Identified system limits & risk areas

- Prioritized remediation roadmap

Outcome

Clear insight into key performance constraints and actionable steps to improve system behavior.

2. Scale & Latency Optimization

Purpose:

Tactical and strategic performance engineering to ensure systems scale reliably under growth and peak demand.

Duration: 4–8 weeks

Investment: \$30,000–\$50,000

Scope

- Throughput and saturation modeling
- Distributed system tuning (JVM, messaging, pipelines)
- Autoscaling and capacity threshold optimization
- SLA / SLO performance hardening
- Load, stress, and scenario validation

Deliverables

- Optimized configuration benchmarks
- Capacity models and future thresholds
- Detailed optimization report
- Executive impact summary

Outcome

Systems tuned for predictable scale with improved throughput, resilience, and reduced latency exposure.

3. Executive Performance Retainer

Purpose:

Ongoing expert oversight to ensure performance, capacity, and scale risk remain managed as platforms evolve.

Duration: Monthly

Investment: \$6,000–\$12,000 per month

Scope

- Monthly performance & capacity reviews
- Forecasting for growth and peak events
- Cloud cost and infrastructure efficiency oversight
- Incident escalation advisory
- Architecture & scale-readiness guidance

Deliverables

- Monthly executive report
- Capacity outlook & risk indicators
- Ongoing optimization recommendations

Outcome

Continuous assurance and strategic guidance to prevent outages, over-provisioning, and inefficient infrastructure use.

4. Incident & Emergency Support

Purpose:

Rapid, on-demand performance diagnostics and escalation support.

Investment: \$350–\$500 per hour

Use Cases

- Production latency spikes
- Capacity failures
- Performance regressions
- High-risk launches or load events

Deliverables

- Immediate performance diagnostics
- Tactical remediation support
- Clear next steps to prevent recurrence

Outcome

Swift system stabilization with priority focus on preventing future incidents.

Approach

KPI99 uses a **structured, evidence-based performance engineering methodology** refined in enterprise settings. We blend deep application insight with infrastructure saturation modeling and capacity forecasting to ensure your systems are reliable, efficient, and predictable.

Why KPI99

- Enterprise-grade performance expertise
 - Outcome-driven service design
 - Executive-level reporting and clarity
 - Senior-led engagements (no junior staff)
 - Clear scope with measurable results
-

Contact

To discuss how KPI99 can support your performance, capacity, and reliability needs:

contact@kpi99.co

<https://kpi99.co/>